



# INFO20003 Database Systems

Assignment Project Exam Help

<https://powcoder.com>

Dr. Renata Borovica-Gajic  
Add WeChat powcoder

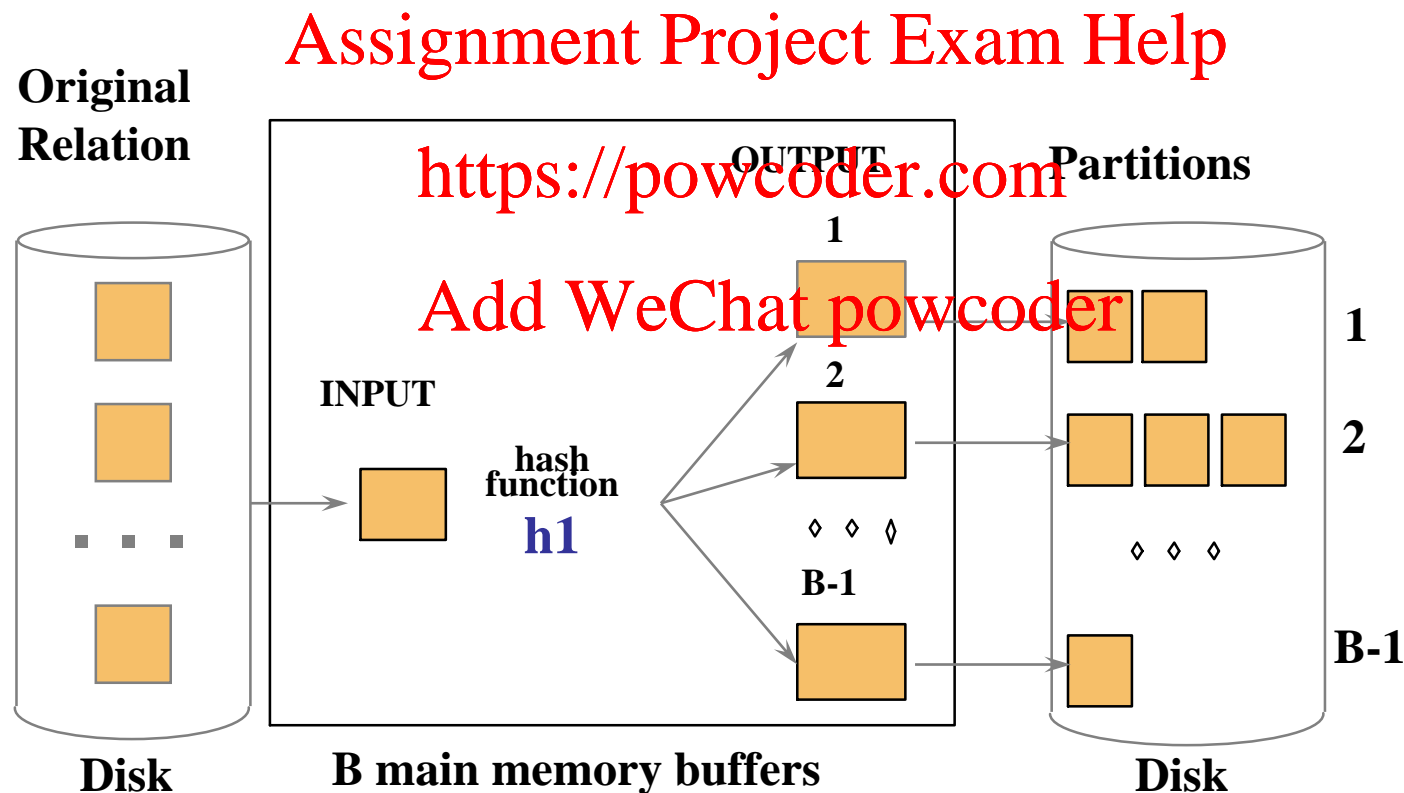
Lecture 12  
Query Processing Part II

Semester 2 2018, Week 6



# Projection based on External Hashing (from last lecture)

1. **Partition** data into B partitions with h1 hash function
2. Load each partition, hash it with another hash function (h2) and **eliminate duplicates**



## 1. Partitioning phase:

- Read R using one input buffer
- For each tuple:
  - Discard unwanted fields
  - Apply hash function  $h_1$  to choose one of B-1 output buffers
- Result is B-1 partitions (of tuples with no unwanted fields)
  - 2 tuples from different partitions guaranteed to be distinct

Assignment Project Exam Help

<https://powcoder.com>

## 2. Duplicate elimination phase:

- For each partition
  - Read it and build an in-memory hash table
    - using hash function  $h_2$  ( $\neq h_1$ ) on all fields
  - while discarding duplicates
- If partition does not fit in memory
  - Apply hash-based projection algorithm recursively to this partition (we will not do this...)

Add WeChat powcoder



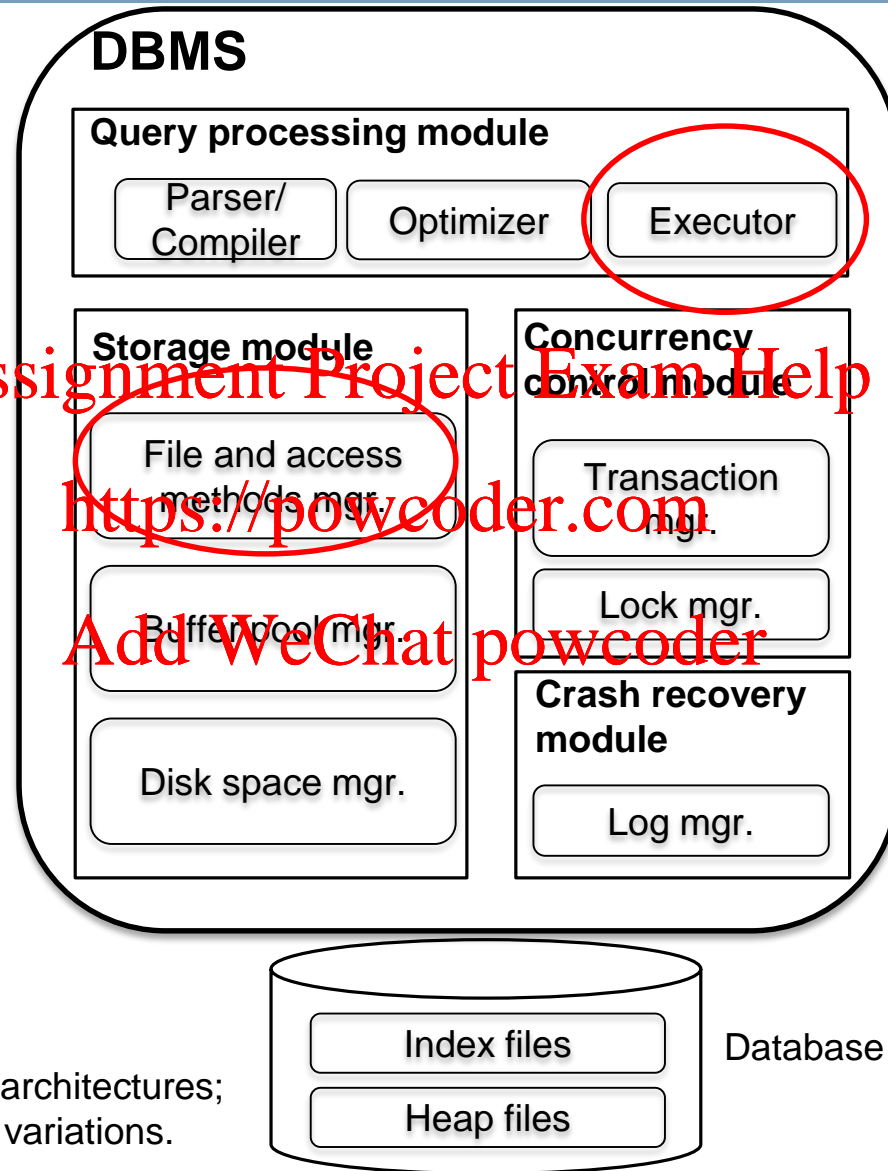
**Cost =** ReadTable + Read the entire table and project attributes  
WriteProjectedPages + Write projected pages into corresponding partitions  
ReadProjectedPages Read partitions one by one, create another hash table and discard duplicates within a bucket

## Assignment Project Exam Help

**Our example:** <https://powcoder.com>

**Cost =** ReadTable + Add WeChat powcoder  
WriteProjectedPages +  
ReadProjectedPages  
=  $1000 + 0.25 * 1000 + 250 = 1500$  (I/O)

# Remember this? Components of a DBMS



**TODAY  
Joins**

**Will briefly  
touch upon ...**

Assignment Project Exam Help

<https://powecoder.com>

Add WeChat powecoder

This is one of several possible architectures; each system has its own slight variations.



- Nested loops join

- Sort-merge join

- Hash join

Assignment Project Exam Help

- General joins

<https://powcoder.com>

Add WeChat powcoder

*Readings: Chapter 14, Ramakrishnan & Gehrke, Database Systems*

- Are very common and can be **very** expensive (cross product in the worst case)
- There are many implementation techniques for join operations **Assignment Project Exam Help**

**<https://powcoder.com>**

- Join techniques we will cover:
  1. Nested-loops join **Add WeChat powcoder**
  2. Sort-merge join
  3. Hash join



**Example:** SELECT \*  
FROM Reserves R1, Sailors S1  
WHERE R1.sid=S1.sid

- In algebra:  $R \bowtie S$ . They are very common and need to be carefully optimized.
- $R \times S$  is large; so,  $R \times S$  followed by a selection is inefficient.

<https://powcoder.com>

Left / Outer	Right / Inner
11,80	11,20
13,75	13,20
12,10	12,20
15,80	13,75
15,44	13,35
13,74	12,10

- Join is associative and commutative:
  - $A \times B == B \times A$
  - $A \times (B \times C) == (A \times B) \times C$
- **Cost metric** : Number of pages; Number of I/O



Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)  
Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)

- **Sailors (S):** Assignment Project Exam Help

- 80 tuples per page, **500 pages**

- $NPages(S) = 500$ ,  $NTuplesPerPage(S) = 80$

- $NTuples(S) = 500 * 80 = 40000$

<https://powcoder.com>  
Add WeChat powcoder

- **Reserves (R):**

- 100 tuples per page, **1000 pages**

- $NPages(R) = 1000$ ,  $NTuplesPerPage(R) = 100$


- $NTuples(R) = 100000$



- For each tuple in the *outer* relation R, we scan the entire *inner* relation S

## Pseudo code:

```
foreach tuple r in R do
  foreach tuple s in S do
    if  $r_i == s_i$  then add  $\langle r, s \rangle$  to result
```



- Cost:

$$\text{Cost (SNJL)} = \text{NPages(Outer)} + \text{NTuples(Outer)} * \text{NPages(Inner)}$$

- Our example:

$$\begin{aligned} \text{Cost (SNLJ)} &= 1000 + 100 * 1000 * 500 \\ &= 50001000 \text{ (I/O)} \end{aligned}$$




- For each **page** of R
  - get each **page** of S
  - write out matching pairs of tuples  $\langle r, s \rangle$ , where  $r$  is in R-page and  $S$  is in S-page

**Pseudo code:** **Assignment Project Exam Help**

```

foreach page  $b_R$  in R do
  foreach page  $b_S$  in S do
    foreach tuple  $r$  in  $b_R$  do
      foreach tuple  $s$  in  $b_S$  do
        if  $r_i == s_j$  then add  $\langle r, s \rangle$  to result
  
```

<https://powcoder.com>  
Add WeChat powcoder

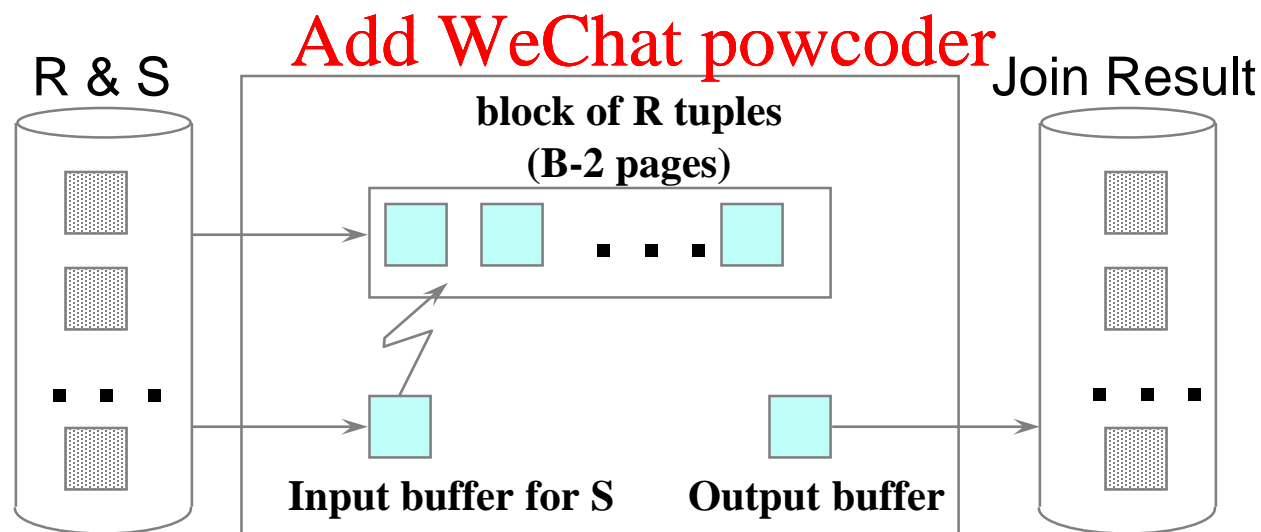


$$\text{Cost (PNJL)} = \text{NPages(Outer)} + \text{NPages(Outer)} * \text{NPages(Inner)}$$

- Our example:

$$\text{Cost (PNLJ)} = 1000 + 1000 * 500 = 501000 \text{ (I/O)}$$

- Page-oriented NL doesn't exploit extra memory buffers
- **Alternative approach:**
  - Use one page as an input buffer for scanning the inner S, one page as the output buffer, and use all remaining pages to hold 'block' of outer R
- For each matching tuple  $r$  in R-block,  $s$  in S-page, add  $\langle r, s \rangle$  to result. Then read next R-block, scan S, etc



$$\text{Cost (BNJL)} = \text{NPages(Outer)} + \text{NBlocks(Outer)} * \text{NPages(Inner)}$$

- $\text{NBlocks(Outer)} = \left\lceil \frac{\text{NPages(Outer)}}{\text{Blocksize} - 1} \right\rceil$

	R	S
B1	11,80	11,20
	13,75	13,20
	12,10	12,20
B2	15,80	13,75
	15,44	13,35
	13,74	12,10

- Our example: <https://powcoder.com>

Let's say we have 102 pages of space in memory, and consider Reserves (R) as the outer and Sailors (S) as the inner table.

$$\text{NBlocks(R)} = 1000 / (102 - 2) = 10$$

$$\text{Cost(BNLJ)} = 1000 + 10 * 500 = 6000 \text{ I/O}$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

• Nested loops join

• Sort-merge join

• Hash join

Assignment Project Exam Help

• General joins

<https://powcoder.com>

Add WeChat powcoder

*Readings: Chapter 14, Ramakrishnan & Gehrke, Database Systems*

# Sort-Merge Join ( $R \bowtie_{i=j} S$ )

- **Sort** R and S on the join column, then scan them to do a **merge** (on join column), and output result tuples

→      R      S      ←

11,80	11,20
12,10	12,20
13,74	13,10
13,73	13,75
15,44	13,35
15,80	13,20

Assignment Project Exam Help  
<https://powcoder.com>

- Sorted R is scanned once:
- Each S group of the same key values is scanned once per matching R tuple (typically means Sorted S is scanned once too).
- Useful when:
  - one or both inputs are already sorted on join attribute(s)
  - output is required to be sorted on join attributes(s)

Add WeChat powcoder



$$\text{Cost (SMJ)} = \text{Sort(Outer)} + \text{Sort(Inner)} \\ + \text{NPages(Outer)} + \text{NPages(Inner)}$$

Sort inputs  
Merge inputs

$$\text{Sort(R)} = \text{External Sort Cost} = 2 * \text{NumPasses} * \text{NPages(R)}$$

Assignment Project Exam Help

<https://powcoder.com>

**Our example:**

Add WeChat powcoder

Let's say that both Reserves and Sailors can be sorted in 2 passes, then:

$$\begin{aligned} \text{Cost(SMJ)} &= \text{Sort R} + \text{Sort S} + \text{NPages(R)} + \text{NPages(S)} \\ &= 2 * 2 * \text{NPages(R)} + 2 * 2 * \text{NPages(S)} \\ &+ \text{NPages(R)} + \text{NPages(S)} \\ &= 5 * 1000 + 5 * 500 = 7500 \text{ I/O} \end{aligned}$$

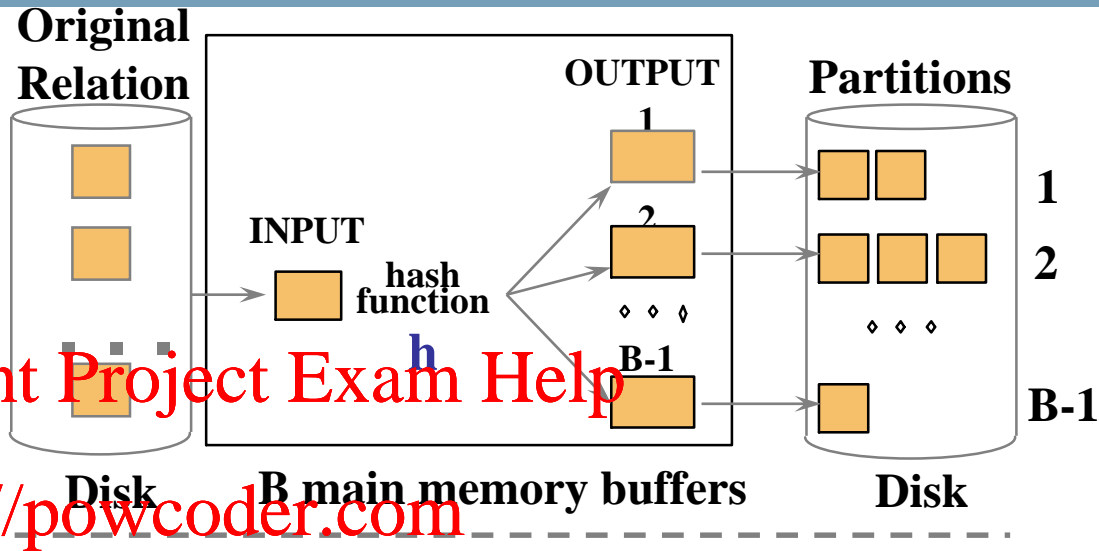


- Nested loops join
  - Sort-merge join
  - Hash join
  - General joins
- Assignment Project Exam Help  
<https://powcoder.com>  
Add WeChat powcoder

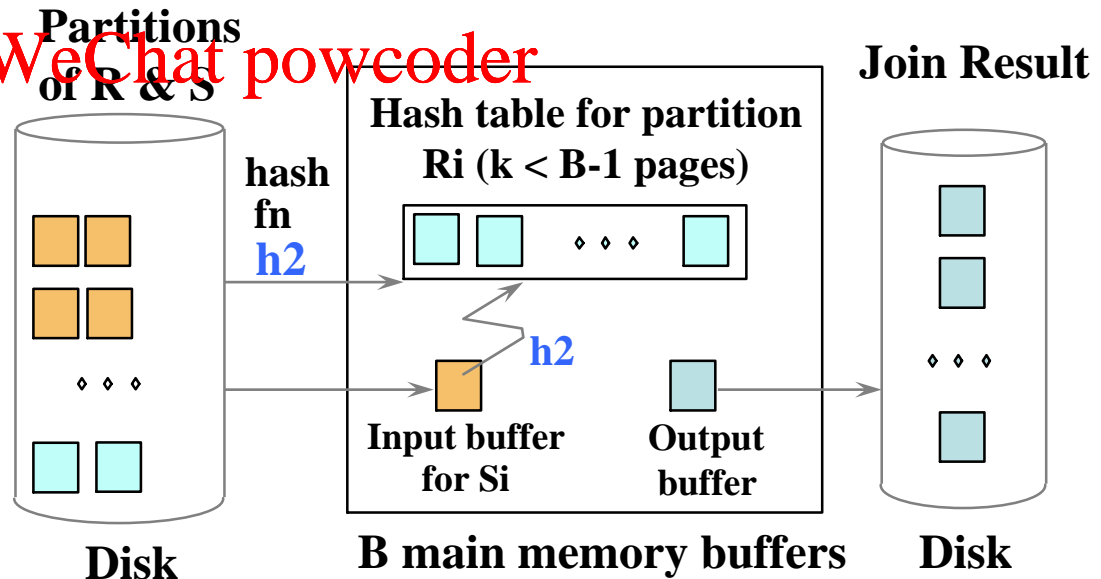
*Readings: Chapter 14, Ramakrishnan & Gehrke, Database Systems*

# Hash-Join

- Partition both relations using hash function  $h$ :  
 $R$  tuples in partition  $l$  will **only** match  $S$  tuples in partition  $l$



- Read in a partition of  $R$ , hash it using  $h_2$  ( $\leftrightarrow h!$ ). Scan matching partition of  $S$ , probe hash table for matches





1. In partitioning phase, we read+write both relations
2. In matching phase, we read both relations

$$\text{Cost (HJ)} = 2 * \text{NPages(Outer)} + 2 * \text{NPages(Inner)} + \text{NPages(Outer)} + \text{NPages(Inner)}$$

Create partitions  
Match partitions

<https://powcoder.com>

• Our example:

Add WeChat powcoder

$$\begin{aligned}\text{Cost(HJ)} &= 2 * \text{NPages(R)} + 2 * \text{NPages(S)} + \text{NPages(R)} + \text{NPages(S)} \\ &= 3 * 1000 + 3 * 500 = 4500 \text{ I/Os}\end{aligned}$$



<https://www.youtube.com/watch?v=o1dMJ6-CKzU>

Assignment Project Exam Help

From 0:58

<https://powcoder.com>

Add WeChat powcoder



# Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



**Start the presentation to activate live content**

If you see this message in presentation mode, install the add-in or get help at [PolleEv.com/app](https://PolleEv.com/app)

## REEDUCATION

- Nested loops join
  - Sort-merge join
  - Hash join
  - General joins
- Assignment Project Exam Help  
<https://powcoder.com>  
Add WeChat powcoder

*Readings: Chapter 14, Ramakrishnan & Gehrke, Database Systems*



- Equalities over several attributes (e.g.,  $R.sid=S.sid$  AND  $R.rname=S.sname$ ):
  - For Sort-Merge and Hash Join, sort/partition on combination of the two join columns
- Inequality conditions (e.g.,  $R.rname < S.sname$ ):
  - Hash Join, Sort Merge Join not applicable
  - Block NL quite likely to be the best join method here

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- A virtue of relational DBMSs:
  - Queries are composed of a few basic operators
  - Implementation of operators can be carefully tuned
  - Important to do this
- Many alternative implementations for each operator
  - No universally superior technique for most operators
- Must consider alternatives for each operation in a query and choose best one based on system statistics...
  - Part of the broader task of optimizing a query composed of several operations

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder





- Understand alternatives for join operator implementations
  - Be able to calculate the cost of alternatives
- Important for Assignment 3 as well

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



- Query optimization
  - How does a DBMS pick a good query plan?

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder