



# INFO20003 Database Systems

Assignment Project Exam Help

<https://powcoder.com>

Dr. Renata Borovica-Gajic  
Add WeChat powcoder

Lecture 13  
Query Optimization Part I

Semester 2 2018, Week 7



- Assignment 1 Feedback will be sent by the end of this week
- Best way to prepare for the MST is to look at your mistakes, compare with the provided solution, and try to mark yourself against the provided assessments criteria
  - Self-reflection is an important part of learning
  - Pay attention to business rules that your model does not support
  - Let's have a look at one potential solution
- This assignment was a preview of a business analyst's life

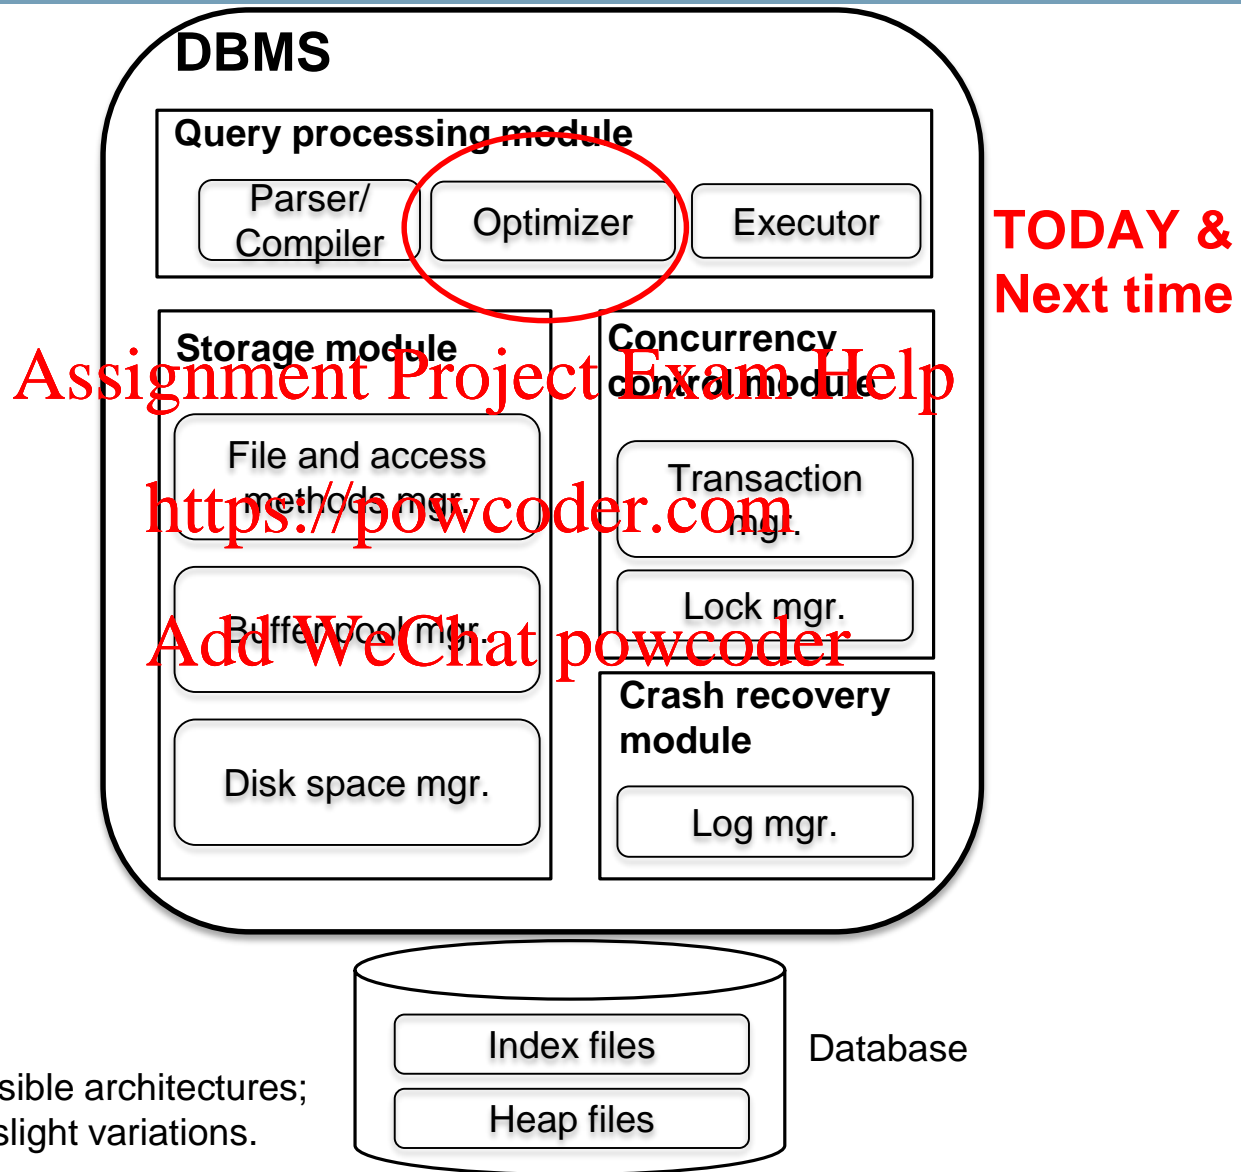
Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



# Remember this? Components of a DBMS



This is one of several possible architectures; each system has its own slight variations.



- Overview

- Query optimization

Assignment Project Exam Help

- Cost estimation

<https://powcoder.com>

Add WeChat powcoder

*Readings: Chapter 12 and 15, Ramakrishnan & Gehrke, Database Systems*



Query

```
Select *  
From Blah B  
Where B.blah = "foo"
```

Query Parser

Assignment Project Exam Help

Query Optimizer

<https://powcoder.com>

Add WeChat powcoder

Plan  
Generator

Plan Cost  
Estimator

Catalog Manager

Schema

Statistics

Query Plan Evaluator



- Typically there are many ways of executing a given query, all giving the same answer
- Cost of alternative methods often **varies enormously**
- Query optimization aims to find the execution strategy with the lowest cost

Assignment Project Exam Help

<https://powcoder.com>

- We will cover:
  - Relational algebra equivalences
  - Cost estimation
  - Result size estimation and reduction factors
  - Enumeration of alternative plans

Add WeChat powcoder

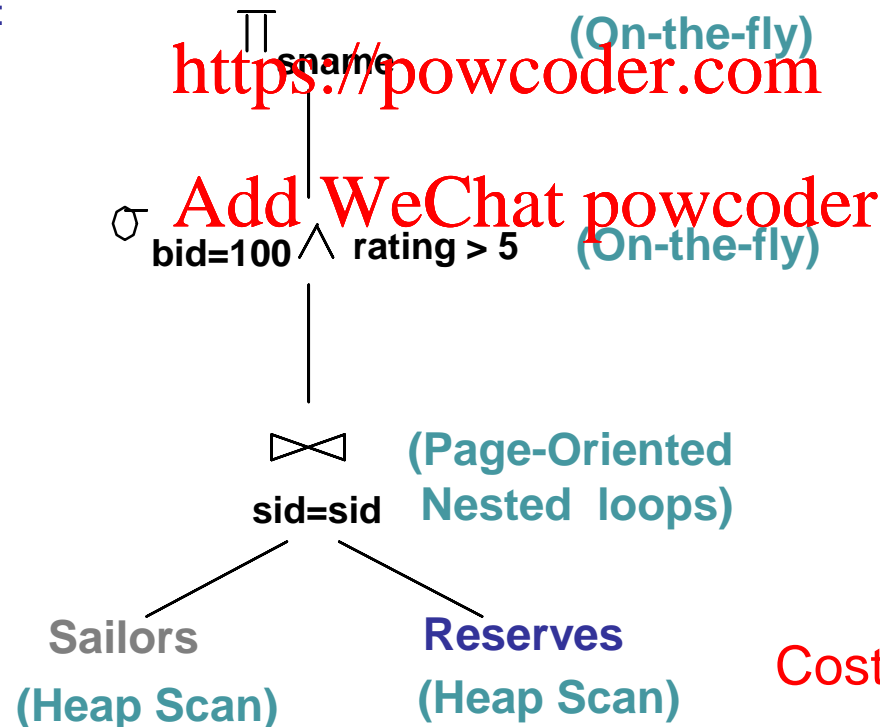


- A tree, with relational algebra operators as nodes
- Each operator labeled with a choice of algorithm

SELECT sname from Sailors NATURAL JOIN Reserves  
WHERE bid = 100 and rating > 5

Assignment Project Exam Help

Plan:



Cost: 500+500\*1000 I/O



- Overview

- Query optimization

Assignment Project Exam Help

- Cost estimation

<https://powcoder.com>

Add WeChat powcoder

*Readings: Chapter 15, Ramakrishnan & Gehrke, Database Systems*





Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)

Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)

Boats (*bid*: integer, *bname*: string, *color*: string)

**Assignment Project Exam Help**

**<https://powcoder.com>**

**Add WeChat powcoder**

Example:

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```

Assignment Project Exam Help

Query optimization steps: <https://powcoder.com>


1. Query first broken into "blocks" [Add WeChat powcoder](#)
2. Each block converted to relational algebra
3. Then, for each block, several alternative query plans are considered
4. Plan with the lowest estimated cost is selected



# Step 1: Break query into query blocks

- Query block is any statement starting with select
- Query block = unit of optimization
- Typically inner most block is optimized first, then moving towards outers

Assignment Project Exam Help



Query:

```
SELECT S.sid  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = "red"
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Relational algebra:

$$\pi_{S.sid}(\sigma_{B.color = \text{"red"}}(Sailors \bowtie Reserves \bowtie Boats))$$



- Selections:  $\sigma_{c_1 \wedge \dots \wedge c_n}(R) \equiv \sigma_{c_1} \left( \dots \left( \sigma_{c_n}(R) \right) \right)$  (Cascade)  
 $\sigma_{c_1} \left( \sigma_{c_2}(R) \right) \equiv \sigma_{c_2} \left( \sigma_{c_1}(R) \right)$  (Commute)
- Projections:  $\pi_{a_1}(R) \equiv \pi_{a_1} \left( \dots \left( \pi_{a_n}(R) \right) \right)$  (Cascade)  
 $a_i$  is a set of attributes of R and  $a_i \subseteq a_{i+1}$  for  $i = 1 \dots n - 1$
- These equivalences allow us to ‘push’ selections and projections ahead of joins.

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



## Selection:

$$\sigma_{\text{age} < 18 \wedge \text{rating} > 5} (\text{Sailors})$$

$$\longleftrightarrow \sigma_{\text{age} < 18} (\sigma_{\text{rating} > 5} (\text{Sailors}))$$

$$\longleftrightarrow \sigma_{\text{rating} > 5} (\sigma_{\text{age} < 18} (\text{Sailors}))$$

## Projection:

~~$$\pi_{\text{age}, \text{rating}} (\text{Sailors}) \longleftrightarrow \pi_{\text{age}} (\pi_{\text{rating}} (\text{Sailors})) \quad ?$$~~

$$\pi_{\text{age}, \text{rating}} (\text{Sailors}) \longleftrightarrow \pi_{\text{age}, \text{rating}} (\pi_{\text{age}, \text{rating}, \text{sid}} (\text{Sailors}))$$



- A projection commutes with a selection that only uses attributes retained by the projection

Assignment Project Exam Help

$\pi_{\text{age, rating, sid}} (\sigma_{\text{age} < 18 \wedge \text{rating} > 5} (\text{Sailors}))$

$\longleftrightarrow \sigma_{\text{age} < 18 \wedge \text{rating} > 5} (\pi_{\text{age, rating, sid}} (\text{Sailors}))$

Add WeChat powcoder

$\pi_{\text{age, sid}} (\sigma_{\text{age} < 18 \wedge \text{rating} > 5} (\text{Sailors}))$

~~$\longleftrightarrow \sigma_{\text{age} < 18 \wedge \text{rating} > 5} (\pi_{\text{age, sid}} (\text{Sailors}))$~~

?

$$R \bowtie (S \bowtie T) \equiv (R \bowtie S) \bowtie T \quad (\text{Associative})$$

$$(R \bowtie S) \equiv (S \bowtie R) \quad (\text{Commutative})$$

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

- These equivalences allow us to choose **different join orders**





- Converting selection + cross-product to join

$$\sigma_{S.sid = R.sid} (\text{Sailors} \times \text{Reserves})$$

$$\leftrightarrow \text{Sailors} \bowtie_{S.sid = R.sid} \text{Reserves}$$

- Selection on just attributes of S commutes with  $R \bowtie S$


$$\sigma_{S.age < 18} (\text{Sailors} \bowtie_{S.sid = R.sid} \text{Reserves})$$

$$\leftrightarrow (\sigma_{S.age < 18} (\text{Sailors})) \bowtie_{S.sid = R.sid} \text{Reserves}$$

- We can also “push down” projection (*but be careful...*)

$$\pi_{S.sname} (\text{Sailors} \bowtie_{S.sid = R.sid} \text{Reserves})$$

$$\leftrightarrow \pi_{S.sname} (\pi_{sname, sid}(\text{Sailors}) \bowtie_{S.sid = R.sid} \pi_{sid}(\text{Reserves}))$$



Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



## REEDUCATION

- Overview
  - Query optimization
  - Cost estimation
- Assignment Project Exam Help
- <https://powcoder.com>
- Add WeChat powcoder

*Readings: Chapter 15, Ramakrishnan & Gehrke, Database Systems*



1. Query first broken into “blocks”
2. Each block converted to relational algebra
3. Then, for each block, several alternative **query plans** are considered
4. Plan with lowest **estimated cost** is selected


**Assignment Project Exam Help**

<https://powcoder.com>

**Add WeChat powcoder**

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```

$\pi_{(sname)} \sigma_{(bid=100 \wedge rating > 5)} (Reserves \bowtie Sailors)$





Queries

```
Select *  
From Blah B  
Where B.blah = "foo"
```

3. What plans are considered?

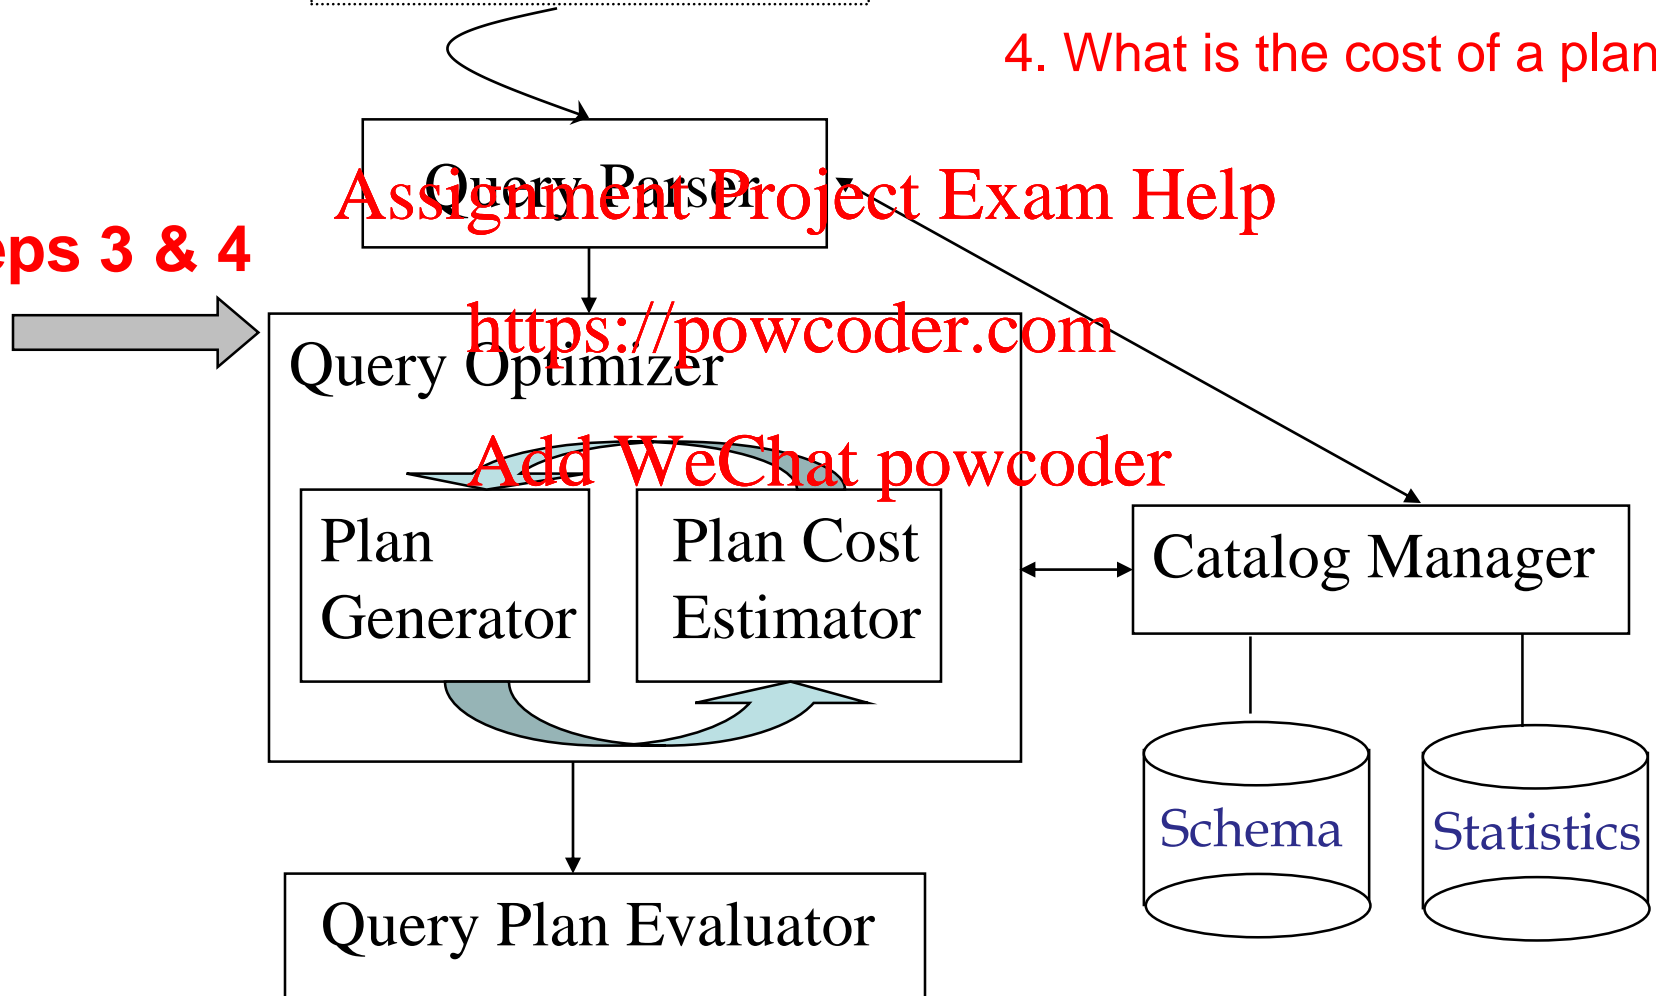
4. What is the cost of a plan?

**Steps 3 & 4**

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder





- For each plan considered, must estimate cost:
  - Must *estimate size of result* for each operation in tree
    - Use information about input relations (from the system catalogs), and apply rules (*discussed next*)
  - Must *estimate cost* of each operation in plan tree
    - Depends on input cardinalities
    - We've already discussed how to estimate the cost of operations (sequential scan, index scan, joins)
    - Next time we will calculate the cost of entire plans...

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



- To decide on the cost, the optimizer needs information about the relations and indexes involved. This information is stored in the system **catalogs**.
- **Catalogs** typically contain at least:
  - # tuples (**NTuples**) and # pages (**NPages**) per relation
  - # distinct key values (**NKeys**) for each index (or relation attribute)
  - low/high key values (**Low/High**) for each index (or relation attribute)
  - Index height (**Height(I)**) for each tree index
  - # index pages (**NPages(I)**) for each index
- Statistics in catalogs are updated periodically

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder



- Consider a query block:

```
SELECT attribute list
FROM relation list
WHERE predicate1 AND ... AND predicate_k
```

Assignment Project Exam Help

- Maximum number of tuples in the result is the **product** of the cardinalities of relations in the FROM clause  
<https://powcoder.com>
- Reduction factor (RF)** associated with each predicate reflects the impact of the predicate in reducing the result size. RF is also called **selectivity**.  
Add WeChat powcoder



- Single table selection:

$$\text{ResultSize} = \text{NTuples}(R) \prod_{i=1..n} RF_i$$

Assignment Project Exam Help

- Joins (over k tables):

<https://powcoder.com>

$$\text{ResultSize} = \prod_{j=1..k} \text{NTuples}(R_j) \prod_{i=1..n} RF_i$$

Add WeChat powcoder

- If there are no selections (no predicates), reduction factors are simply ignored, i.e. they are ==1





- Depend on the type of the predicate:
  1. Col = value  
**RF = 1/NKeys(Col)**
  2. Col > value  
**RF = (High(Col) – value) / (High(Col) – Low(Col))**  
<https://powcoder.com>
  3. Col < value  
**RF = (val – Low(Col)) / (High(Col) – Low(Col))**  
Add WeChat powcoder
  4. Col\_A = Col\_B (for joins)  
**RF = 1/ (Max (NKeys(Col\_A), NKeys(Col\_B)))**
  5. In no information about Nkeys or interval, use a “magic number” 1/10  
**RF = 1/10**



Sailors (S): NTuples(S) = 1000, Nkeys(rating) = 10 interval [1-10],  
age interval [0-100], Nkeys(sid)=1000

*SELECT \* FROM Sailors WHERE rating = 3 AND age > 50;*

**Assignment Project Exam Help**  
**Calculate result size:**

<https://powcoder.com>

**NTuples(S) = 1000**

**RF(rating) =  $1/10 = 0.1$**

**RF(age) =  $(100-50)/(100-0) = 0.5$**

**ResultSize = NTuples(S)\*RF(rating)\*RF(age)**

**=  $1000*0.1*0.5 = 50$  tuples**

**Add WeChat powcoder**

**Given Reserves (R): NTuples(R) = 100, Nkeys(sid) = 100 and  
Sailors (S): NTuples(S) = 1000, Nkeys(rating) = 10 interval [1-  
10], age interval [0-100], Nkeys(sid) = 1000 and the query:  
SELECT \* FROM Sailors as S INNER JOIN Reserves as R ON S.SID  
= R.SID WHERE rating = 8 and 20 < age < 30; Calculate Result Size**

**Assignment Project Exam Help**

**<https://powcoder.com>**

**Add WeChat powcoder**



- What is query optimization/describe steps?
- Equivalence classes
- Result size estimation

## Assignment Project Exam Help

- Important for Assignment 3 as well

<https://powcoder.com>

Add WeChat powcoder



- Query optimization Part II
  - Plan enumeration

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder